

Shubham Bhattacharya

AI Agent Backend Engineer — TypeScript, Node.js, LangGraph.js

shubhambhattacharya107@gmail.com | +91-9155252394 | Vadodara, Gujarat, India
github.com/shubhambhattacharya-dev | linkedin.com/in/shubhambhattadev | shubhambhattacharya.dev

SUMMARY

AI Agent Backend Engineer specializing in production LLM systems, multi-agent orchestration, and intelligent automation with TypeScript and Node.js. Built agentic backends using LangGraph.js state machines with deterministic routing, Zod-validated outputs, and human-in-the-loop patterns. Strong focus on zero-trust AI architecture where every LLM output is untrusted until validated. Seeking AI Agent Backend Developer roles building production-grade agent systems.

TECHNICAL SKILLS

AI Agent Development: LangGraph.js, Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, Prompt Engineering
LLM Integration: Multi-Provider LLM Orchestration, LLM Output Validation, tiktoken, Streaming (SSE)
Retrieval & Embeddings: pgvector, RAG Pipelines, Semantic Search, BM25, Vector Databases
Languages: TypeScript, JavaScript, SQL
Backend & APIs: Node.js, Express.js, Next.js, REST APIs, SSE, Socket.IO, WebSockets
Database & Cache: PostgreSQL, Prisma, Redis, BullMQ, MongoDB
Validation & Auth: Zod, Clerk, JWT, bcrypt, Helmet
Observability: Langfuse, OpenTelemetry, Sentry, Pino
DevOps & Infra: Docker, GitHub Actions, Vercel

PROJECTS

Gigi AI Ops — Production AI Agent Backend

LangGraph.js, TypeScript, Zod, pgvector, PostgreSQL, BullMQ, Redis

- Architected multi-step AI agent backend for e-commerce operations using LangGraph.js state machines — deterministic intent classification, tool calling, and graceful failure recovery
- Built multi-provider LLM orchestration layer with automatic fallback across Groq, Gemini, and OpenRouter — per-persona routing with cost and latency tradeoffs
- Implemented hybrid RAG pipeline combining BM25 keyword search and dense vector search via pgvector for zero-hallucination FAQ retrieval
- Enforced zero-crash reliability via Zod runtime validation on every LLM output — unstructured AI responses never reach the database
- Integrated tiktoken WASM for per-request token usage tracking and cost calculation in USD/INR
- Designed human-in-the-loop approval flow for financial operations — refunds above threshold require admin review

github.com/shubhambhattacharya-dev/Agentic_Project | agentic-project-drab.vercel.app

Medo Copilot — Multi-Provider AI Launch Auditor

Next.js, TypeScript, Vercel AI SDK, Playwright, Clerk, Tailwind CSS

- Built AI-powered launch auditor with multi-provider LLM waterfall fallback: Gemini 2.0, Groq/Llama 4 Scout, OpenRouter/Claude
- Designed hybrid scoring engine: 60% AI insights + 40% deterministic metrics from Lighthouse and static code analysis
- Implemented Playwright-based visual intelligence capturing desktop and mobile screenshots for multimodal AI review
- Built static security analyzer scanning for hardcoded secrets and code quality issues with structured JSON output

github.com/shubhambhattacharya-dev/Medo_Copilot | medo-copilot.shubhambhattacharya.dev

DocNow — Telemedicine SaaS Platform

Next.js, TypeScript, Clerk, Prisma, PostgreSQL, Vonage, shadcn/ui

- Built full-stack telemedicine platform with WebRTC video calls via Vonage API, role-based dashboards, and credit-based payments
- Designed credit-based payment system with full audit trail: purchases, deductions, refunds, and earnings with transaction history
- Implemented doctor payout system with platform fee calculation and admin approval workflow

github.com/shubhambhattacharya-dev/health | docnow-zeta.vercel.app

EDUCATION

BCA — Parul University | 2022 — 2025 | CGPA: 8.0/10

Relevant Coursework: Data Structures & Algorithms, DBMS, Computer Networks, Operating Systems, Web Technologies